

IOPS Scaling on Azure Database for PostgreSQL Flexible Server

Prerequisites:

1. You must be under NDA.
2. IOPS Scaling feature can be enabled for newly created or existing server servers during private preview.

Limitations:

1. Private preview is only available for the customers who are under DNA.

IOPS Scaling Feature – Scaling Up

Follow these steps to scale up the Performance Tier.

1. In the [Azure Portal](#), choose your flexible server that you want to scale up.
2. Click **Compute + storage**.
3. A page with the current settings is displayed.

The screenshot shows the 'Settings' page for an Azure Database for PostgreSQL Flexible Server. The 'Compute + storage' tab is selected and highlighted with a red box. The settings are as follows:

Setting	Value
Compute Processor	Intel
Compute size	Standard_D2ds_v4 (2 vCores, 8 GiB memory, 3200 max iops)
Storage	
The storage you provision is the amount of storage capacity available to your flexible server and is billed GiB/month. Note that storage cannot be scaled down once the server is created.	
Storage size	128 GiB
Performance Tier (preview)	P10 (500 iops)
Storage Auto-growth (preview)	<input type="checkbox"/>

4. You will see the new "Performance Tier (preview)" drop-down option. The option selected will be the pre-provisioned IOPS, which is also the minimum amount of IOPS available.

Settings

- Compute + storage
- Networking
- Databases
- Connection strings
- Server parameters
- Replication
- Maintenance
- High availability
- Backup and restore
- Advisor recommendations
- Locks

Power Platform

- Power BI (preview)

Security

Compute Processor

☐ AMD

☒ Intel

Compute size

Standard_D2ds_v4 (2 vCores, 8 GiB memory, 3200 max iops)

Storage

The storage you provision is the amount of storage capacity available to your flexible server and is billed GiB/month.
Note that storage cannot be scaled down once the server is created.

Storage size

128 GiB

Performance Tier (preview)

P10 (500 iops)

P10 (500 iops)

P15 (1100 iops)

P20 (2300 iops)

P30 (5000 iops)

P40 (7500 iops)

P50 (7500 iops)

5. Select your new performance tier and click save.

Settings

- Compute + storage
- Networking
- Databases
- Connection strings
- Server parameters
- Replication
- Maintenance
- High availability
- Backup and restore
- Advisor recommendations

Compute

Compute resources are pre-allocated and billed per hour based on vCores configured.
Note that high availability is supported for only General purpose and Memory optimized tiers.

Compute tier

☐ Burstable (1-20 vCores) - Best for workloads that don't need the full CPU continuously

☒ General Purpose (2-96 vCores) - Balanced configuration for most common workloads

☐ Memory Optimized (2-96 vCores) - Best for workloads that require a high memory to CPU ratio

Compute Processor

☐ AMD

☒ Intel

Compute size

Standard_D2ds_v4 (2 vCores, 8 GiB memory, 3200 max iops)

Storage

The storage you provision is the amount of storage capacity available to your flexible server and is billed GiB/month.
Note that storage cannot be scaled down once the server is created.

Storage size

128 GiB

Performance Tier (preview)

P30 (5000 iops)

Storage Auto-growth (preview)

☐

6. Your server will deploy and once the deployment is completed, your server is updated and will show the new performance tier.

IOPS Scaling Feature – Scaling Down

Follow these steps to scale down the Performance Tier.

1. In the [Azure Portal](#), choose your flexible server that you want to scale down.
2. Click **Compute + storage**.
3. A page with the current settings is displayed.

Settings

- Compute + storage
- Networking
- Databases
- Connection strings
- Server parameters
- Replication
- Maintenance
- High availability
- Backup and restore
- Advisor recommendations

Compute Processor

Compute size

Standard_D2ds_v4 (2 vCores, 8 GiB memory, 3200 max iops)

Storage

The storage you provision is the amount of storage capacity available to your flexible server and is billed GiB/month.
Note that storage cannot be scaled down once the server is created.

Storage size ①

128 GiB

Performance Tier (preview) ①

P30 (5000 iops)

Storage Auto-growth (preview) ①

☐

4. You will see the new “Performance Tier (preview)” drop-down option. The option selected will be your last selected IOPS when you scaled up.

Settings

- Compute + storage
- Networking
- Databases
- Connection strings
- Server parameters
- Replication
- Maintenance
- High availability
- Backup and restore
- Advisor recommendations
- Locks
- Power Platform
- Power BI (preview)
- Security

Compute Processor

Compute size

Standard_D2ds_v4 (2 vCores, 8 GiB memory, 3200 max iops)

Storage

The storage you provision is the amount of storage capacity available to your flexible server and is billed GiB/month.
Note that storage cannot be scaled down once the server is created.

Storage size ①

128 GiB

Performance Tier (preview) ①

P30 (5000 iops)

Storage Auto-growth (preview) ①

High availability

High availability ①

Backups

Backup retention period (in days)

5. Select your new performance tier and click save.

The screenshot shows the AWS Management Console configuration page for a new EC2 instance. The left sidebar contains navigation links: Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, Migration, Settings, Compute + storage (selected), Networking, Databases, Connection strings, Server parameters, Replication, Maintenance, High availability, and Backup and restore. The top navigation bar includes a search bar, a back arrow, and buttons for Save (highlighted with a red box), Discard, CLI / PS, Feedback, and FAQs. The main content area is titled 'Compute' and contains the following sections:

- Compute tier:** Three radio button options:
 - Burstable (1-20 vCores) - Best for workloads that don't need the full CPU continuously
 - General Purpose (2-96 vCores) - Balanced configuration for most common workloads** (selected)
 - Memory Optimized (2-96 vCores) - Best for workloads that require a high memory to CPU ratio
- Compute Processor:** Two radio button options:
 - AMD
 - Intel** (selected)
- Compute size:** A dropdown menu showing 'Standard_D2ds_v4 (2 vCores, 8 GiB memory, 3200 max iops)'.
- Storage:** A section with a note: 'The storage you provision is the amount of storage capacity available to your flexible server and is billed GiB/month. **Note that storage cannot be scaled down once the server is created.**'
- Storage size:** A dropdown menu showing '128 GiB'.
- Performance Tier (preview):** A dropdown menu (highlighted with a red box) showing 'P15 (1100 iops)' selected.

6. Your server will deploy and once the deployment is completed, your server is updated and will show the new performance tier.